

```

from google.colab import drive

drive.mount('/content/gdrive')

Mounted at /content/gdrive

cd /content/gdrive/My Drive/k means clustering

/content/gdrive/My Drive/k means clustering

ls

'case study 1(1).pdf'      gender_classification_v7.csv  README.md
'case study 1.pdf'        k-means-clustering.ipynb
'Copy of driver-data.csv' LICENSE

import pandas as pd, numpy as np, matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression as LR
from sklearn.neighbors import KNeighborsClassifier as KNC
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB as GNB
from sklearn.tree import DecisionTreeClassifier as DTC
from sklearn.ensemble import RandomForestClassifier as RFC
import seaborn as s
from sklearn.metrics import accuracy_score

```

▼ Importing the Dataset -

```

#loading the dataset
data=pd.read_csv("gender_classification_v7.csv")

```

```
data.head() #print first five rows
```

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	di
0	1	11.8	6.1	1	0	1	
1	0	14.0	5.4	0	0	1	
2	0	11.8	6.3	1	1	1	
3	0	14.4	6.1	0	1	1	
4	1	13.5	5.9	0	0	0	

```
data.tail() #print last five rows
```

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	di
4996	1	13.6	5.1	0	0	0	
4997	1	11.9	5.4	0	0	0	
4998	1	12.9	5.7	0	0	0	
4999	1	13.2	6.2	0	0	0	
5000	1	15.4	5.4	1	1	1	

▼ EDA

```

#statistical measure of the data
data.describe()

```

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	distance_nose_to_lip_long	
count	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000	
mean	0.869626	13.181484	5.946311	0.493901	0.507898	0.493101	0.498900	
std	0.336748	1.107128	0.541268	0.500013	0.499988	0.500002	0.500049	
min	0.000000	11.400000	5.100000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	12.200000	5.500000	0.000000	0.000000	0.000000	0.000000	
50%	1.000000	13.100000	5.900000	0.000000	1.000000	0.000000	0.000000	

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5001 entries, 0 to 5000
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   long_hair                            5001 non-null  int64
1   forehead_width_cm                   5001 non-null  float64
2   forehead_height_cm                  5001 non-null  float64
3   nose_wide                           5001 non-null  int64
4   nose_long                           5001 non-null  int64
5   lips_thin                           5001 non-null  int64
6   distance_nose_to_lip_long           5001 non-null  int64
7   gender                              5001 non-null  object
dtypes: float64(2), int64(5), object(1)
memory usage: 312.7+ KB

print(np.shape(data))

(5001, 8)
```

The above analysis tells us that the dataset has 14 columns and 303 columns and the data types are "int64" and "float64"

```
print(data.columns) #prints name of the columns

Index(['long_hair', 'forehead_width_cm', 'forehead_height_cm', 'nose_wide',
       'nose_long', 'lips_thin', 'distance_nose_to_lip_long', 'gender'],
      dtype='object')

print(data.isnull().any()) # checking for any null value in the data set

long_hair      False
forehead_width_cm  False
forehead_height_cm False
nose_wide      False
nose_long      False
lips_thin      False
distance_nose_to_lip_long False
gender         False
dtype: bool

data= data.replace(['Male','Female'],[1, 0])

data.head()
```

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	distance_nose_to_lip_long	gender
0	1	11.8	6.1	1	0	1	1	1
1	0	14.0	5.4	0	0	1	0	0
2	0	11.8	6.3	1	1	1	1	1
3	0	14.4	6.1	0	1	1	1	1
4	1	13.5	5.9	0	0	0	0	0

result shows there are no null values

```
#checking on the target variable
data["gender"].value_counts()

0    2501
1    2500
Name: gender, dtype: int64
```

```
X=data.drop(["gender"],axis=1)
X
```

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	distance_nose_to_lip_long
0	1	11.8	6.1	1	0	1	1
1	0	14.0	5.4	0	0	1	0
2	0	11.8	6.3	1	1	1	1
3	0	14.4	6.1	0	1	1	1
4	1	13.5	5.9	0	0	0	0
...
4996	1	13.6	5.1	0	0	0	0
4997	1	11.9	5.4	0	0	0	0
4998	1	12.9	5.7	0	0	0	0
4999	1	13.2	6.2	0	0	0	0
5000	1	15.4	5.4	1	1	1	1

5001 rows × 7 columns

```
y=data.gender
y
```

```
0    1
1    0
2    1
3    1
4    0
..
4996  0
4997  0
4998  0
4999  0
5000  1
Name: gender, Length: 5001, dtype: int64
```

▼ Splitting data into train_set and test_set

```
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=0)
```

```
#checking the splitted value
```

```
print(x_test.shape,x_train.shape,y_test.shape,y_train.shape)
```

```
(1001, 7) (4000, 7) (1001,) (4000,)
```

performing some model and training the data set

▼ Logistic Regression

```
#choosing a model currently logistic regression
```

```
log=LR()
log.fit(x_train,y_train) #training a model x_train.T to transpose
```

```
print("Test accuracy of training data using Logistic regression: {}".format(log.score(x_train,y_train)*100))
print("Test accuracy of test data using Logistic regression: {}".format(log.score(x_test,y_test)*100))
```

```
Test accuracy of training data using Logistic regression: 96.925
Test accuracy of test data using Logistic regression: 96.5034965034965
```

▼ K nearest neighbour

```
#k nearest neighbour
```

```
knn=KNC(n_neighbors=7)
knn.fit(x_train,y_train)
```


```
print("Test accuracy of training data using KNN: {}".format(knn.score(x_train,y_train)*100))
print("Test accuracy of test data using KNN regression: {}".format(knn.score(x_test,y_test)*100))
```

```
Test accuracy of training data using KNN: 97.575
Test accuracy of test data using KNN regression: 96.2037962037962
```


```
# to get the best accuracy
score=[]
for i in range(1,20):
    knn=KNC(n_neighbors=i)
    knn.fit(x_train,y_train)
    prediction=knn.predict(x_test)
    score.append(knn.score(x_test,y_test)*100)
print("the best accuracy is",max(score),"at",score.index(max(score))+1)
```

```
plt.plot(range(1,20),score
)
plt.xlabel("K value")
plt.ylabel("Score")
plt.show()
```

the best accuracy is 97.002997002997 at 12



```
from sklearn.metrics import confusion_matrix,classification_report
cm=confusion_matrix(y_test,knn_predicting)
s.heatmap(cm,annot=True)
```



Support Vector machine



support vector machine algorithm

```
svm=SVC(random_state=1)
svm.fit(x_train,y_train)
print("Test accuracy of training data using SVM: {}".format(svm.score(x_train,y_train)*100))
print("Test accuracy of test data using SVM regression: {}".format(svm.score(x_test,y_test)*100))
```

Test accuracy of training data using SVM: 97.35000000000001
Test accuracy of test data using SVM regression: 96.6033966033966

Naive Bayes Algorithm

#naive bayes algo

```
nb=GNB()
nb.fit(x_train,y_train)
print("Test accuracy of training data using NB: {}".format(nb.score(x_train,y_train)*100))
print("Test accuracy of test data using NB regression: {}".format(nb.score(x_test,y_test)*100))
```

Test accuracy of training data using NB: 97.275
Test accuracy of test data using NB regression: 96.8031968031968

Decision Tree Algorithm

decision tree algo

```
dtc=DTC()
dtc.fit(x_train,y_train)
print("Test accuracy of training data using DTC regression: {}".format(dtc.score(x_train,y_train)*100))
print("Test accuracy of test data using DTC regression: {}".format(dtc.score(x_test,y_test)*100))
```

Test accuracy of training data using DTC regression: 99.925
Test accuracy of test data using DTC regression: 96.00399600399601

Random Forest algorithm

random forest algorithm

```
rf=RFC(n_estimators=1000,random_state=1)
rf.fit(x_train,y_train)
print("Test accuracy of training data using RFC regression: {}".format(rf.score(x_train,y_train)*100))
print("Test accuracy of test data using RFC regression: {}".format(rf.score(x_test,y_test)*100))
```


Test accuracy of training data using RFC regression: 99.925
Test accuracy of test data using RFC regression: 96.30369630369631

```
methods=["LR","KNN","SVM","NB","DTC","RFC"]
accuracy=[log.score(x_test,y_test)*100,knn.score(x_test,y_test)*100,svm.score(x_test,y_test)*100,nb.score(x_test,y_test)*100,dtc.score(x_test,y_test)*100,rf.score(x_test,y_test)*100]
colors=["purple","green","blue","orange","magenta","#CFC60E","#0FBBAE"]
```

```
s.set_style("whitegrid")
plt.figure(figsize=(16,8))
```

```
plt.ylabel("accuracy",size=20)  
plt.xlabel("methods")  
s.barplot(x=methods,y=accuracy,palette=colors)
```

<Axes: xlabel='methods', ylabel='accuracy'>



▼ Applying KNN to predict the output

```
knn_predicting=knn.predict(x_test)  
print("y true :",y_test[:10].values)  
print("y predict :",knn_predicting[:10])
```

```
y true : [0 1 1 0 0 0 0 1 0 0]  
y predict : [0 1 1 0 0 0 0 1 0 0]
```